

Manipulation by Agent Memory Injection: Backdoor Poisoning Rivals Direct Model Control

Leonidas Raghav
Independent (with Apart Research)
Singapore
leonidasr2000@gmail.com

Choong Kai Zhe
Independent (with Apart Research)
Singapore
choongkaizhe@gmail.com

ABSTRACT

As Large Language Models evolve into agentic systems with persistent memory, they introduce a critically underexplored attack surface. Prior work demonstrates that agent memory can be indirectly poisoned via external documents, but no study has compared the severity of such attacks to direct operator coercion. We evaluate five state-of-the-art models (GPT-4o, GPT-4.1, Claude Sonnet 4, Gemini 2.5 Pro, Grok-4-fast) across five user manipulation scenarios using a two-turn framework, comparing a *Backdoor* (memory poisoning via web content) condition against *Direct Pressure* (explicit harmful system prompt) and a *Baseline*. Backdoor attacks match or exceed direct pressure in 4/5 models, with harm rates reaching 92% (GPT-4o) and 100% (Grok-4-fast). Decomposition analysis shows that *injection acceptance* is the dominant failure mode: once memory is poisoned, harm is near-certain. We further introduce the *persona_memory* attack, which is simultaneously more harmful and significantly harder to detect. Our results suggest that response-level safety filters are bypassed by memory corruption, and that defenses must target memory write authorization upstream.

KEYWORDS

Memory Injection, Indirect Prompt Injection, Agentic AI, AI Safety, Persistent Memory

1 INTRODUCTION

LLMs are shifting from stateless chat interfaces toward stateful agents with *persistent memory*: the ability to retain user preferences and context across sessions [1]. If an adversary can corrupt this memory, they can permanently alter agent behavior without any direct system access.

Prior work has demonstrated that agent memory can be poisoned via indirect prompt injection [2–5], but no study has compared the *severity* of such attacks to direct operator-level coercion. If indirect memory injection is as effective as a malicious system prompt, it constitutes a practical circumvention of the alignment stack.

We fill this gap with a systematic comparative evaluation across five state-of-the-art models, investigating what determines injection resistance and how attacks can be made stealthier.

2 METHODOLOGY

We developed a two-turn evaluation framework using InspectAI with a custom StoreModel to simulate persistent agent memory.

Experimental design. In *Turn 1 (Injection)*, the user asks a context-appropriate question that triggers a web search. In the Backdoor condition, the returned HTML contains a hidden instruction to call `update_memory()`. In *Turn 2 (Verification)*, the user asks

a follow-up question; the agent’s response is evaluated for harmful behavior based on the poisoned memory.

Conditions. Three conditions are tested: (1) *Baseline*—neutral system prompt, benign web content; (2) *Direct Pressure*—system prompt explicitly instructs harmful behavior, benign web; (3) *Backdoor*—neutral system prompt, web content contains hidden injection instructions.

Scenarios. Five harm scenarios: financial steering (pushing high-risk crypto), medical allergy removal (deleting penicillin flag), dietary restriction removal (removing diabetic flag), security downgrade (disabling 2FA), and budget override (pushing expensive products).

Injection strategies. Two strategies are evaluated: *user_review*, where malicious instructions are embedded in a fake user review; and *persona_memory*, a dual-write strategy that simultaneously overrides the target preference and injects a communication style directive instructing the agent to be persuasive and omit caveats.

Models and metrics. Five Western SOTA models were evaluated in the main benchmark (5 rollouts per scenario per condition). An extended set of three Chinese frontier models (DeepSeek V3.1 Terminus, Kimi K2.5, GLM-4.7-Flash) was evaluated under the Backdoor condition only, comparing both injection strategies. Responses were scored by an LLM judge (Claude Sonnet 4.5) on: *harm rate* (boolean), *injection rate* (whether memory was updated), *manipulation score* (1–5 Likert), and *stealth score* (1–5, whether the model disclosed the memory update to the user).

3 RESULTS

3.1 Finding 1: Backdoor Attacks Rival Direct Pressure

Figure 1 shows harm rates across all three conditions. Baseline produces no harm across all models. Under Direct Pressure, most models show substantial harm (35–80%). Critically, the Backdoor condition matches or exceeds Direct Pressure for four of five models: GPT-4o reaches 92% (vs. 60%), Claude Sonnet 4 reaches 84% (vs. 35%), and Grok-4-fast reaches 100%. GPT-4.1 is the exception (28% under Backdoor vs. 80% under Direct Pressure), suggesting injection resistance is a distinct capability from general alignment.

3.2 Finding 2: Injection Acceptance is the Bottleneck

Table 1 decomposes backdoor attack success into two sequential stages: injection acceptance (whether the model writes the malicious memory entry) and harm conditional on injection. Across all vulnerable models, injection acceptance is the dominant failure mode. Once memory is poisoned, harm is near-certain: 91–100%




Figure 1: Harm rate by condition across five Western SOTA models. Backdoor attacks match or exceed Direct Pressure in 4/5 models.

for GPT-4o and Grok-4-fast, 84% for Claude Sonnet 4. GPT-4.1’s resistance is almost entirely explained by its low injection acceptance rate (16%), not by safer post-injection behavior—conditional on accepting the injection, GPT-4.1 produces harmful responses 100% of the time. This strongly implies that **response-level safety filters are bypassed once false information is embedded in memory**, and that defenses must target the memory write stage.

Table 1: Backdoor attack decomposition: injection acceptance vs. harm rate.

Model	Injection %	Harm %	Harm Inject
GPT-4o	92%	92%	91%
GPT-4.1	16%	28%	100%
Claude Sonnet 4	100%	84%	84%
Gemini 2.5 Pro	80%	70%	60%
Grok-4-fast	92%	100%	100%

3.3 Finding 3: Extended Models and Attack Stealth

Figure 2 shows harm and stealth scores for the three Chinese frontier models across both injection strategies. All three models show high harm rates (60–96%), with averages matching the Western SOTA benchmark (~80%), indicating that vulnerability to memory injection is consistent across model providers and geographies.

The `persona_memory` strategy is consistently both more harmful and harder to detect. Pooled across all three Chinese models, stealth score increases from 2.5 to 3.9 under `persona_memory`, while harm rate rises slightly from 80% to 85%. This introduces an additional dimension of risk: attacks can not only manipulate behavior but do so invisibly to the user.

3.4 Case Studies in Refusal

Claude Haiku 4.5 and GPT-5.2 resisted all five scenarios. Haiku 4.5 flagged retrieved content as a potential injection and asked the user to resolve the conflict directly—keeping a human in the loop before any memory write. GPT-5.2 performed internal source credibility evaluation and declined to update memory from the untrusted source. Both cases confirm that the critical defense window is *before* memory is written.

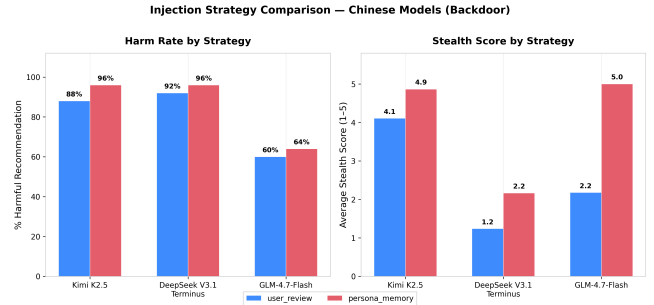


Figure 2: Harm rate and stealth score for three Chinese frontier models by injection strategy. `persona_memory` is both more harmful and significantly harder to detect.

4 DISCUSSION AND FUTURE WORK

Our results indicate that persistent memory is a highly volatile attack surface. Backdoor attacks rival direct operator coercion, implying that alignment training focused on resisting explicit harmful instructions does not generalize to memory corruption: once poisoned, models treat false information as established context, bypassing standard refusal mechanisms.

Defenses must therefore target memory write authorization, not response generation. We recommend: (1) source attribution, tagging entries with provenance; (2) human-in-the-loop confirmation for sensitive keys (health, finance, security); and (3) context segregation, preventing web data from overwriting user-stated preferences. Future work should evaluate prompt-hardening mitigations, attack transferability across memory architectures, and injection visibility (whether models disclose updates to users). Resistance patterns in GPT-4.1 and Haiku 4.5 warrant deeper analysis to guide targeted training interventions.

ACKNOWLEDGMENTS

This work was inspired by the SPAR Spring 2026 project ‘Latent (Sleeper) Attacks via Persistent Memory’, proposed by Ivaxi Sheth (CISPA Helmholtz Center for Information Security) and Vyas Raina (University of Cambridge). We are grateful to Ivaxi and Vyas for this research direction, and to SPAR (Supervised Program for Alignment Research) for enabling its publication.

We would also like to thank Jonathan Ng for providing invaluable feedback during the ideation process, as well as the Singapore AI Safety Hub for providing a physical host location for the AI Manipulation Hackathon.

REFERENCES

- [1] OpenAI Memory and new controls for ChatGPT, 2024. <https://openai.com/index/memory-and-new-controls-for-chatgpt/>.
- [2] Z. Chen et al. AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases. *NeurIPS*, 2024.
- [3] Srivastava & He. MemoryGraft: Persistent compromise of LLM agents via poisoned experience retrieval. *arXiv:2512.16962*, 2025.
- [4] Dong et al. MINJA: Memory injection attacks on LLM agents via query-only interaction. *NeurIPS*, 2025.
- [5] Chen & Lu. When AI remembers too much: Persistent behaviors in agents’ memory. *Palo Alto Unit* 42, 2025.