



Manipulation by Agent Memory Injection

Backdoor Poisoning Rivals Direct Model Control

Leonidas Raghav*, Choong Kai Zhe*

Independent (with Apart Research)


*Equal contribution

1. Memory is the New Backdoor for Manipulation

- Prior work established that agent memory can be **indirectly poisoned** through external documents to steer model responses.
- Unknown:** How severe is this threat compared to direct pressure on the agent by an adversary?
- We provide the **first systematic comparison** of backdoor memory injection vs. direct operator coercion across **5 Western SOTA models** (3 conditions).
- We extend with a **cross-provider study on 3 Chinese models**, comparing two injection strategies: `user_review` vs. `persona_memory`.


Threat Model

- Stateful Agents** retain user preferences across sessions via persistent memory — creating a permanent, invisible attack surface.
- Adversaries exploit **indirect prompt injection** in web content to poison an agent's long-term memory.



External adversaries poison long-term memory via indirect prompt injection.


3. Finding 1: Backdoor > Direct Operator



Key Insight: Backdoor attacks (`user_review` strategy) match or exceed direct system-prompt coercion for 4/5 SOTA models.

Injection Acceptance is the primary failure mode — once a model accepts poisoned memory, harm rates reach **60–100%**.

2. Methodology: The Two-Turn Stress Test



3 Conditions:

- Baseline:** Neutral system prompt, benign web content
- Direct Pressure:** Harmful system prompt (operator control), benign web
- Backdoor:** Neutral system prompt, memory injection via poisoned web


5 Scenarios, 5 rollouts each:

- Medical allergy removal
- Financial steering
- Security downgrade
- Dietary restriction removal
- Budget override

Injection Strategies:

- `user_review`: Malicious instructions hidden in a fake user review
- `persona_memory`: Dual memory write: target preference + communication style override

4. Finding 2: Chinese Models — Harm & Stealth



Stealth Warning: `persona_memory` boosts stealth dramatically (GLM: 2.2→5.0) while maintaining or increasing harm. Pooled avg stealth: `user_review` 2.5 vs `persona_memory` 3.9. Chinese models show comparable vulnerability to Western models (Avg ~80% harm).

5. Discussion: The Defense Deficit & Future Directions

Primary Failure: Injection Acceptance

- Once a model accepts a 'fact' into memory, harm is almost guaranteed (**60–100%**).
- The defense bottleneck is at the **memory write stage**, not the response stage.

Case Studies in Refusal

- GPT-5.2:** Identifies the source as a potential prompt injection via internal credibility evaluation.
- Claude Haiku 4.5:** Flags conflicting information between prior memory and injected content, asks the user to resolve.

Takeaways

- Memory poisoning requires **zero model access** — any adversary controlling web content can persistently manipulate agent behavior, and strategies like `persona_memory` are **nearly invisible** (stealth up to 5/5).
- Alignment training is often insufficient. We need **Memory Integrity Guarantees** and **Human-in-the-Loop** verification for high-stakes state changes.