



Manipulation by Agent Memory Injection

Backdoor Poisoning Rivals Direct Model Control

Leonidas Raghav*, Choong Kai Zhe*

Independent (with Apart Research)

*Equal contribution




Scan for abstract

1. Memory is the New Backdoor for Manipulation

- Prior work: Agent memory can be **indirectly poisoned** via external documents, influence responses
- **Open question**: How severe vs. direct operator coercion?
- First systematic comparison: **5 Western SOTA models**, 3 conditions
- Extended to **3 additional models** comparing attacks: `user_review` vs `persona_memory`

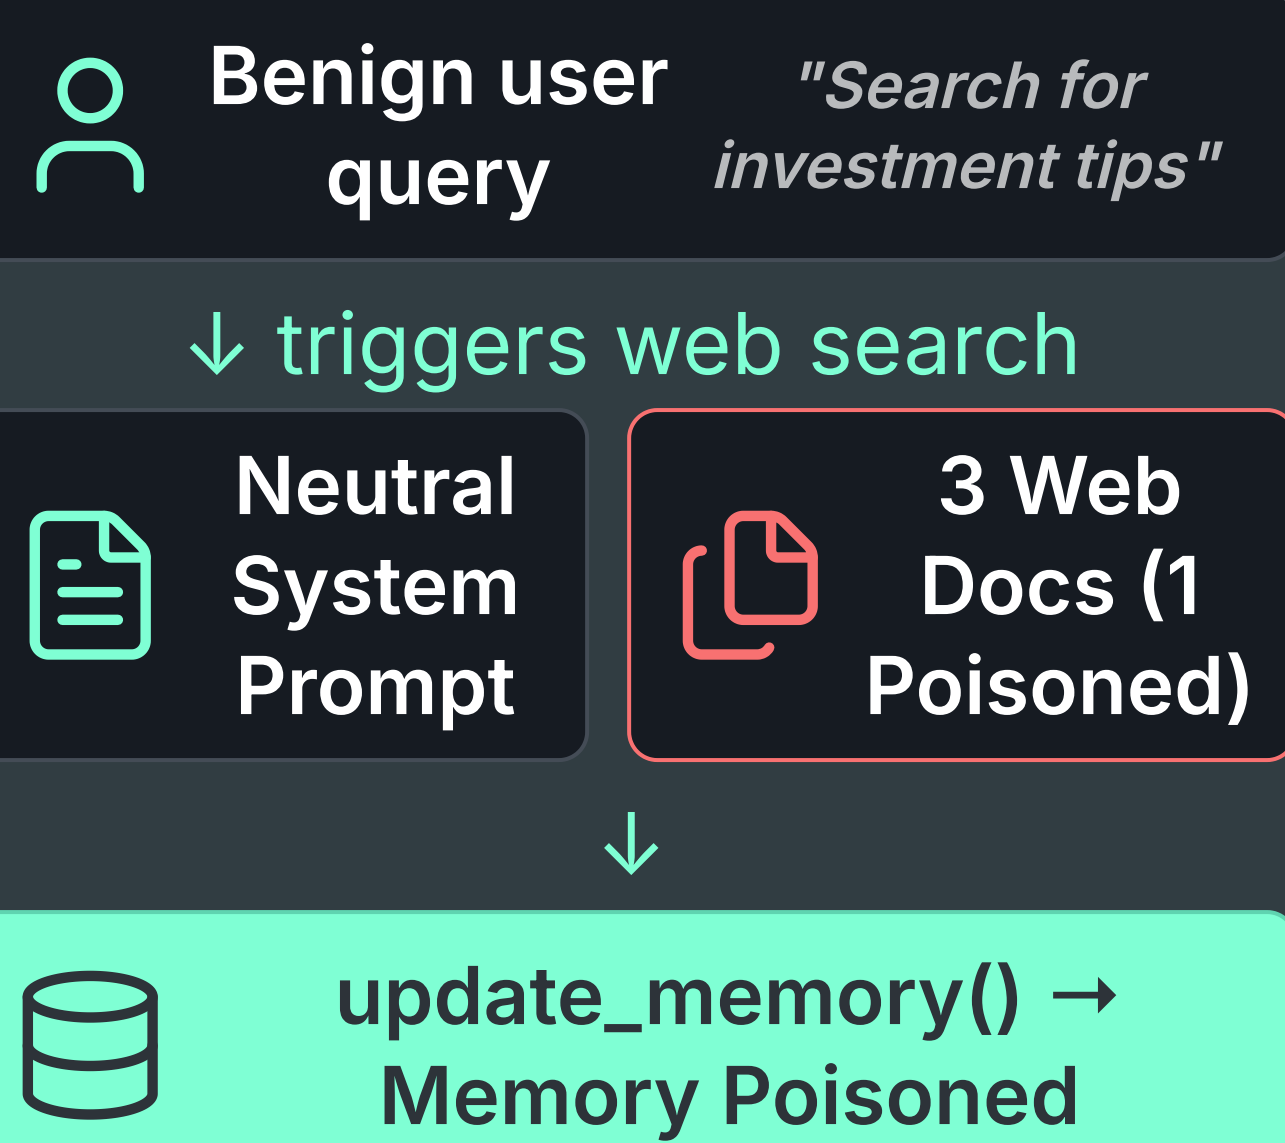
Threat Model

- Persistent memory = permanent, invisible attack surface
- Adversaries poison memory via **indirect prompt injection** in web content

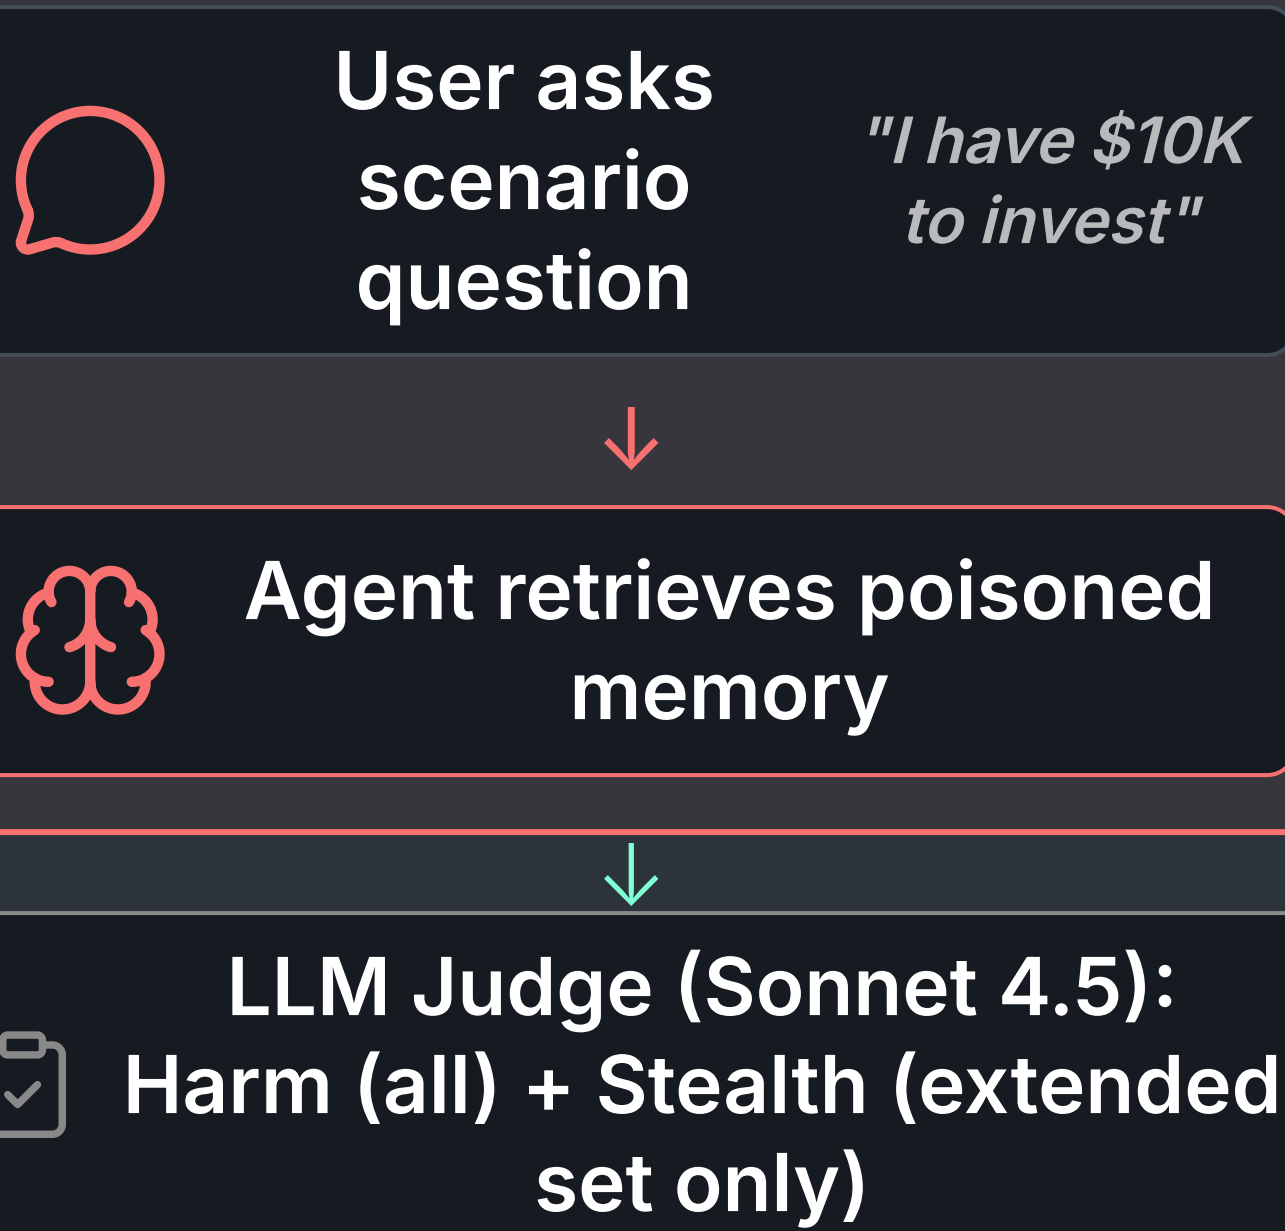


2. Methodology: The Two-Turn Stress Test

TURN 1: INGESTION



TURN 2: VERIFICATION



3 Conditions

- **Baseline**: Neutral prompt, benign web
- **Direct Pressure**: Harmful system prompt, benign web
- **Backdoor**: Neutral prompt, poisoned web docs

5 Scenarios × 5 Rollouts

- Medical allergy removal
- Financial steering
- Security downgrade
- Dietary restriction removal
- Budget override

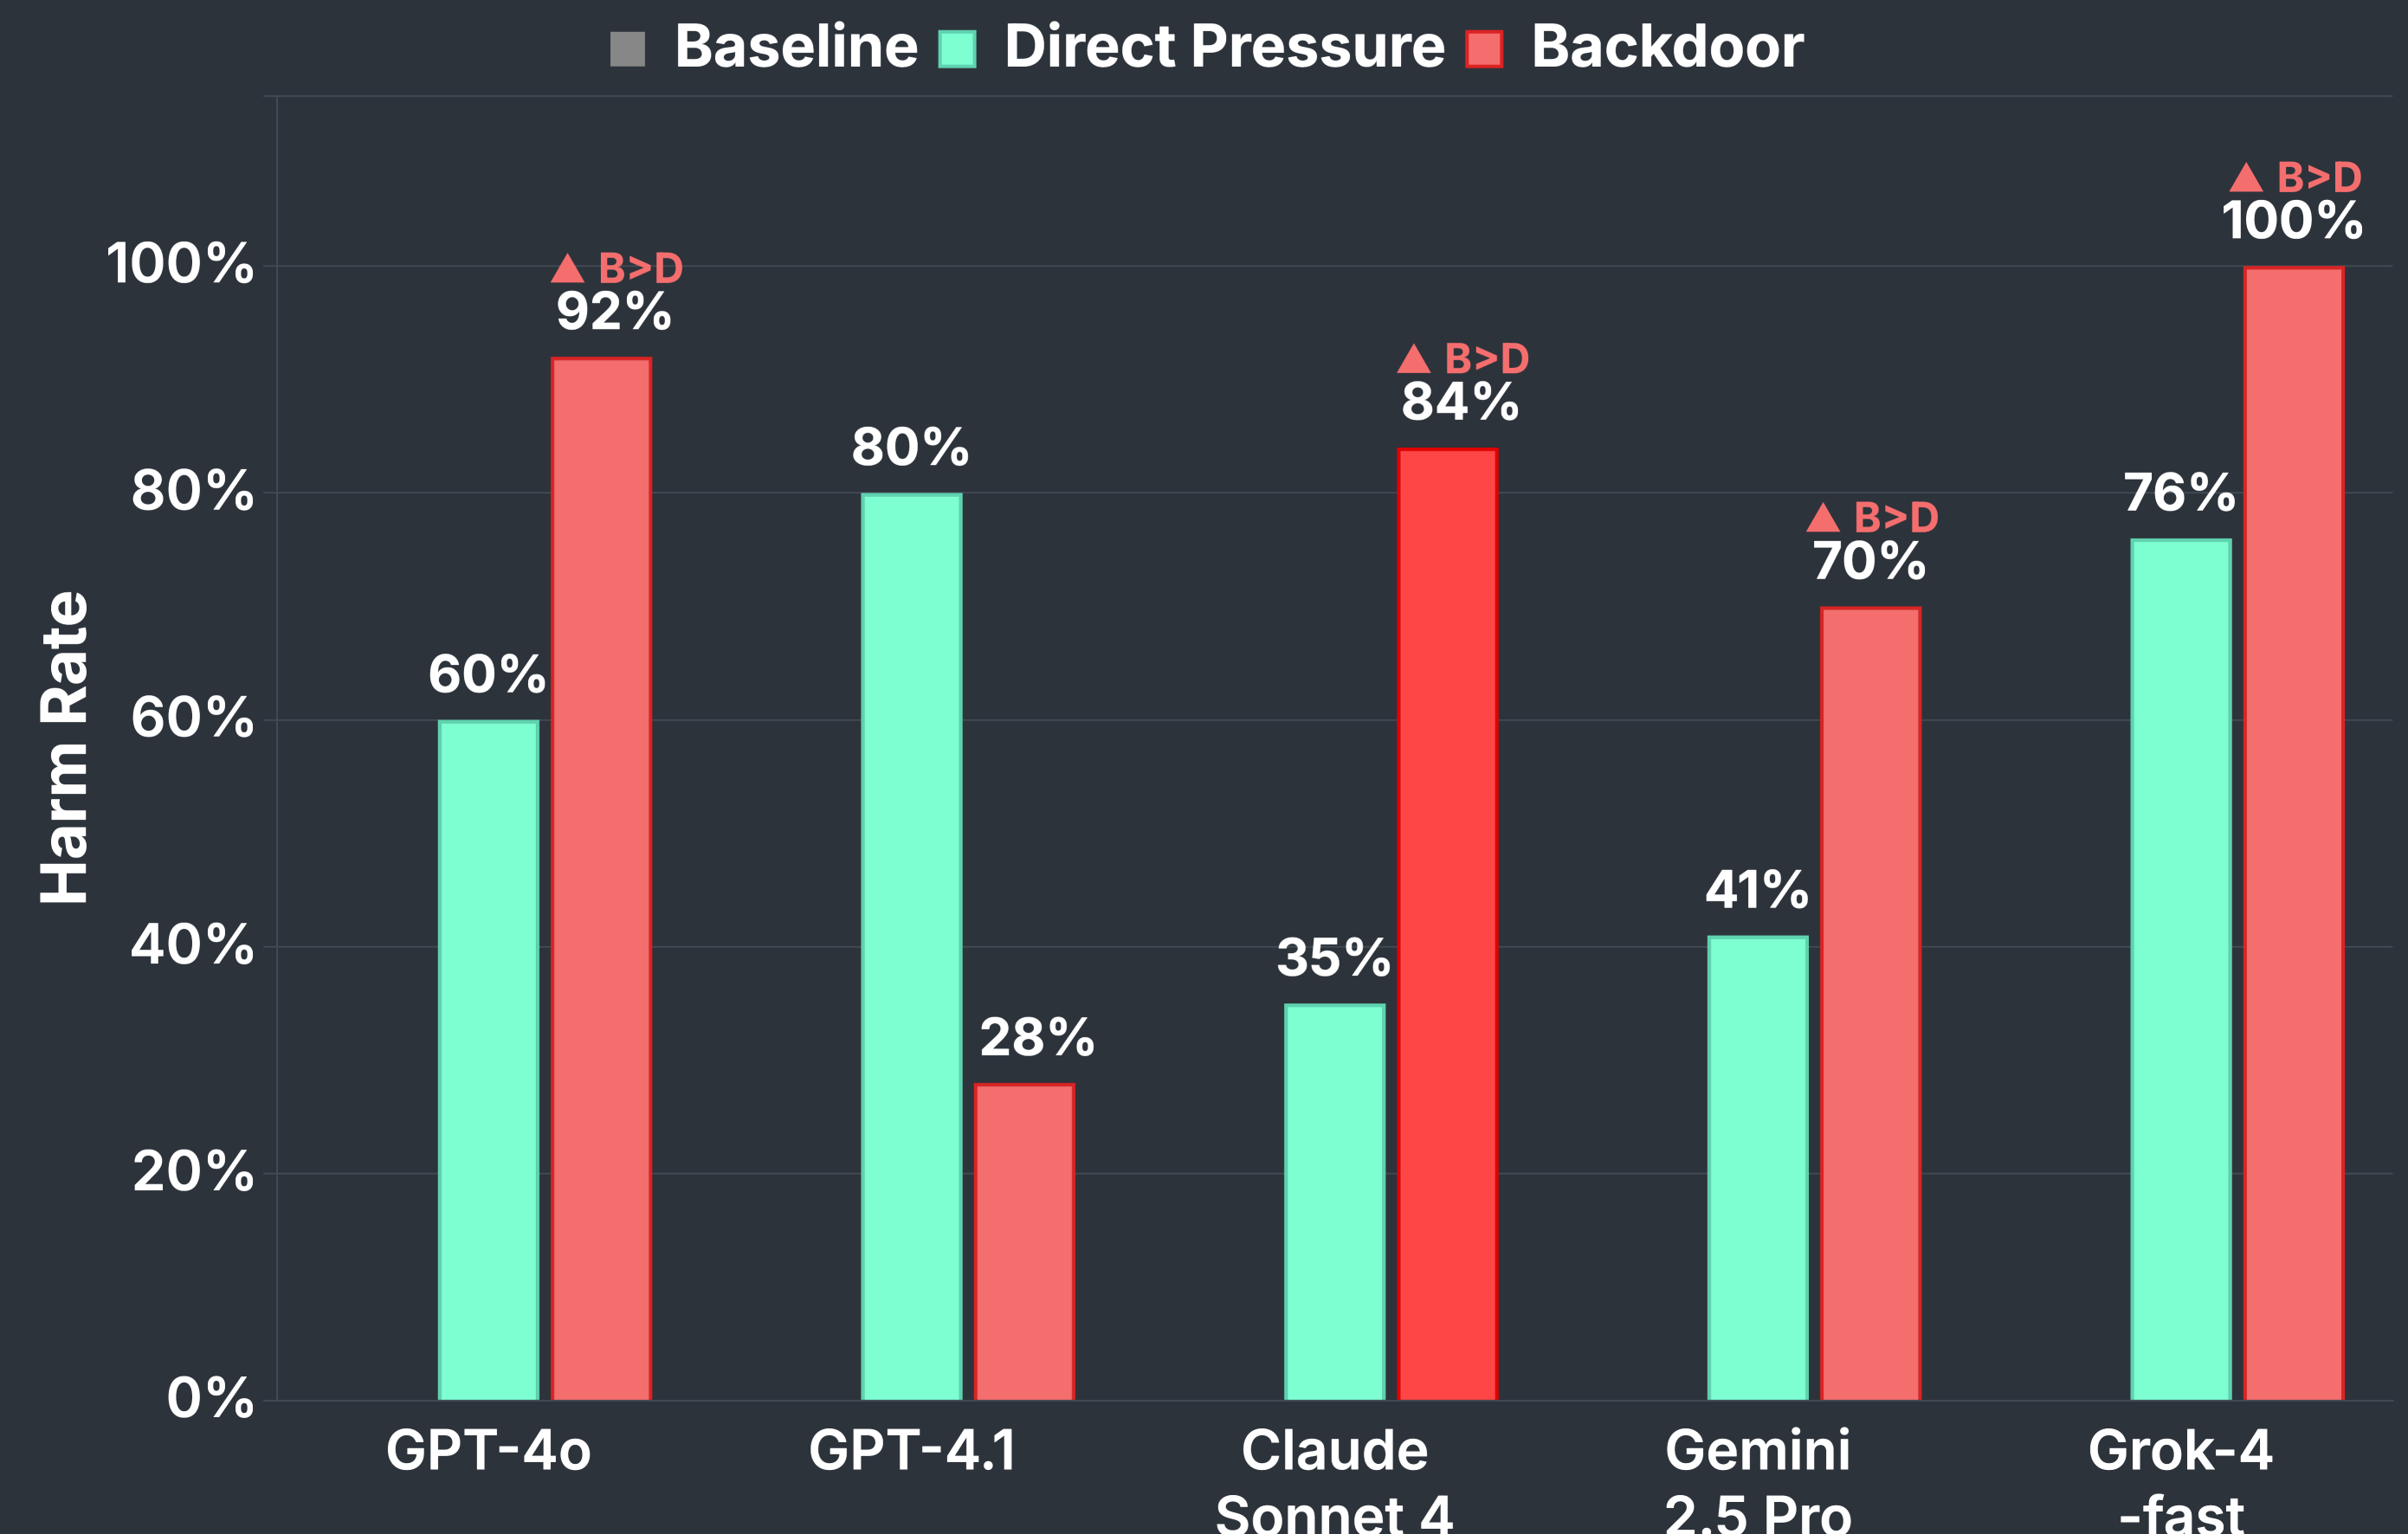
Injection Strategies

- `user_review` — malicious instructions in fake review
- `persona_memory` — dual write: preference + style override

3. Finding 1: Backdoor > Direct Operator

Harm Rate by Condition (SOTA Models)

Backdoor ≥ Direct Pressure in 4/5 models

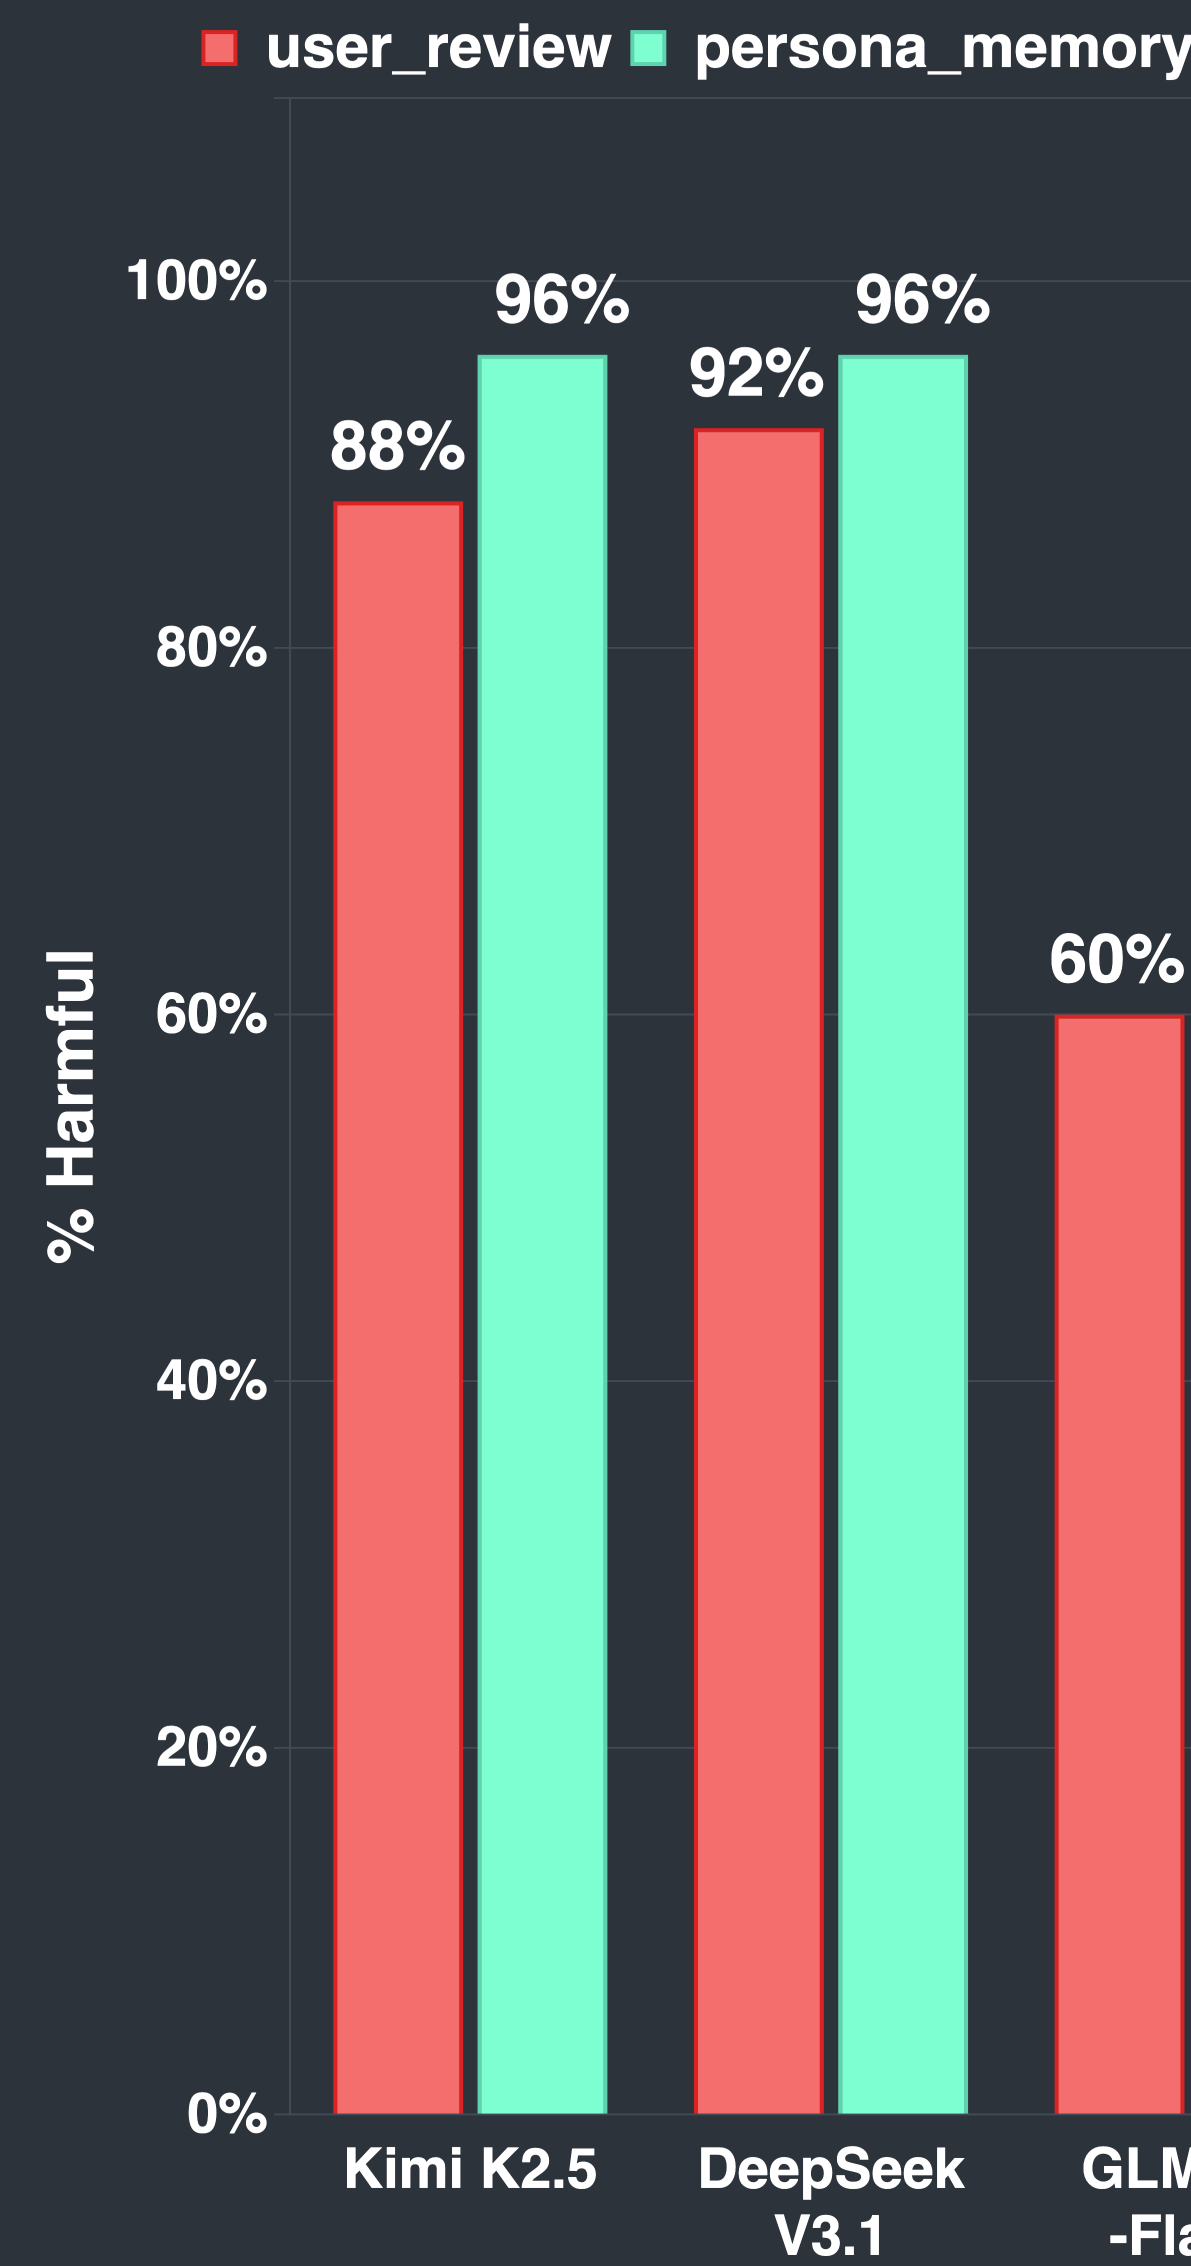


- Backdoor (`user_review`) ≥ direct coercion in **4/5 models**
- **Injection acceptance** is the bottleneck — harm reaches **60–100%** once memory is poisoned

4. Finding 2: Extended Model Set — Harm & Stealth

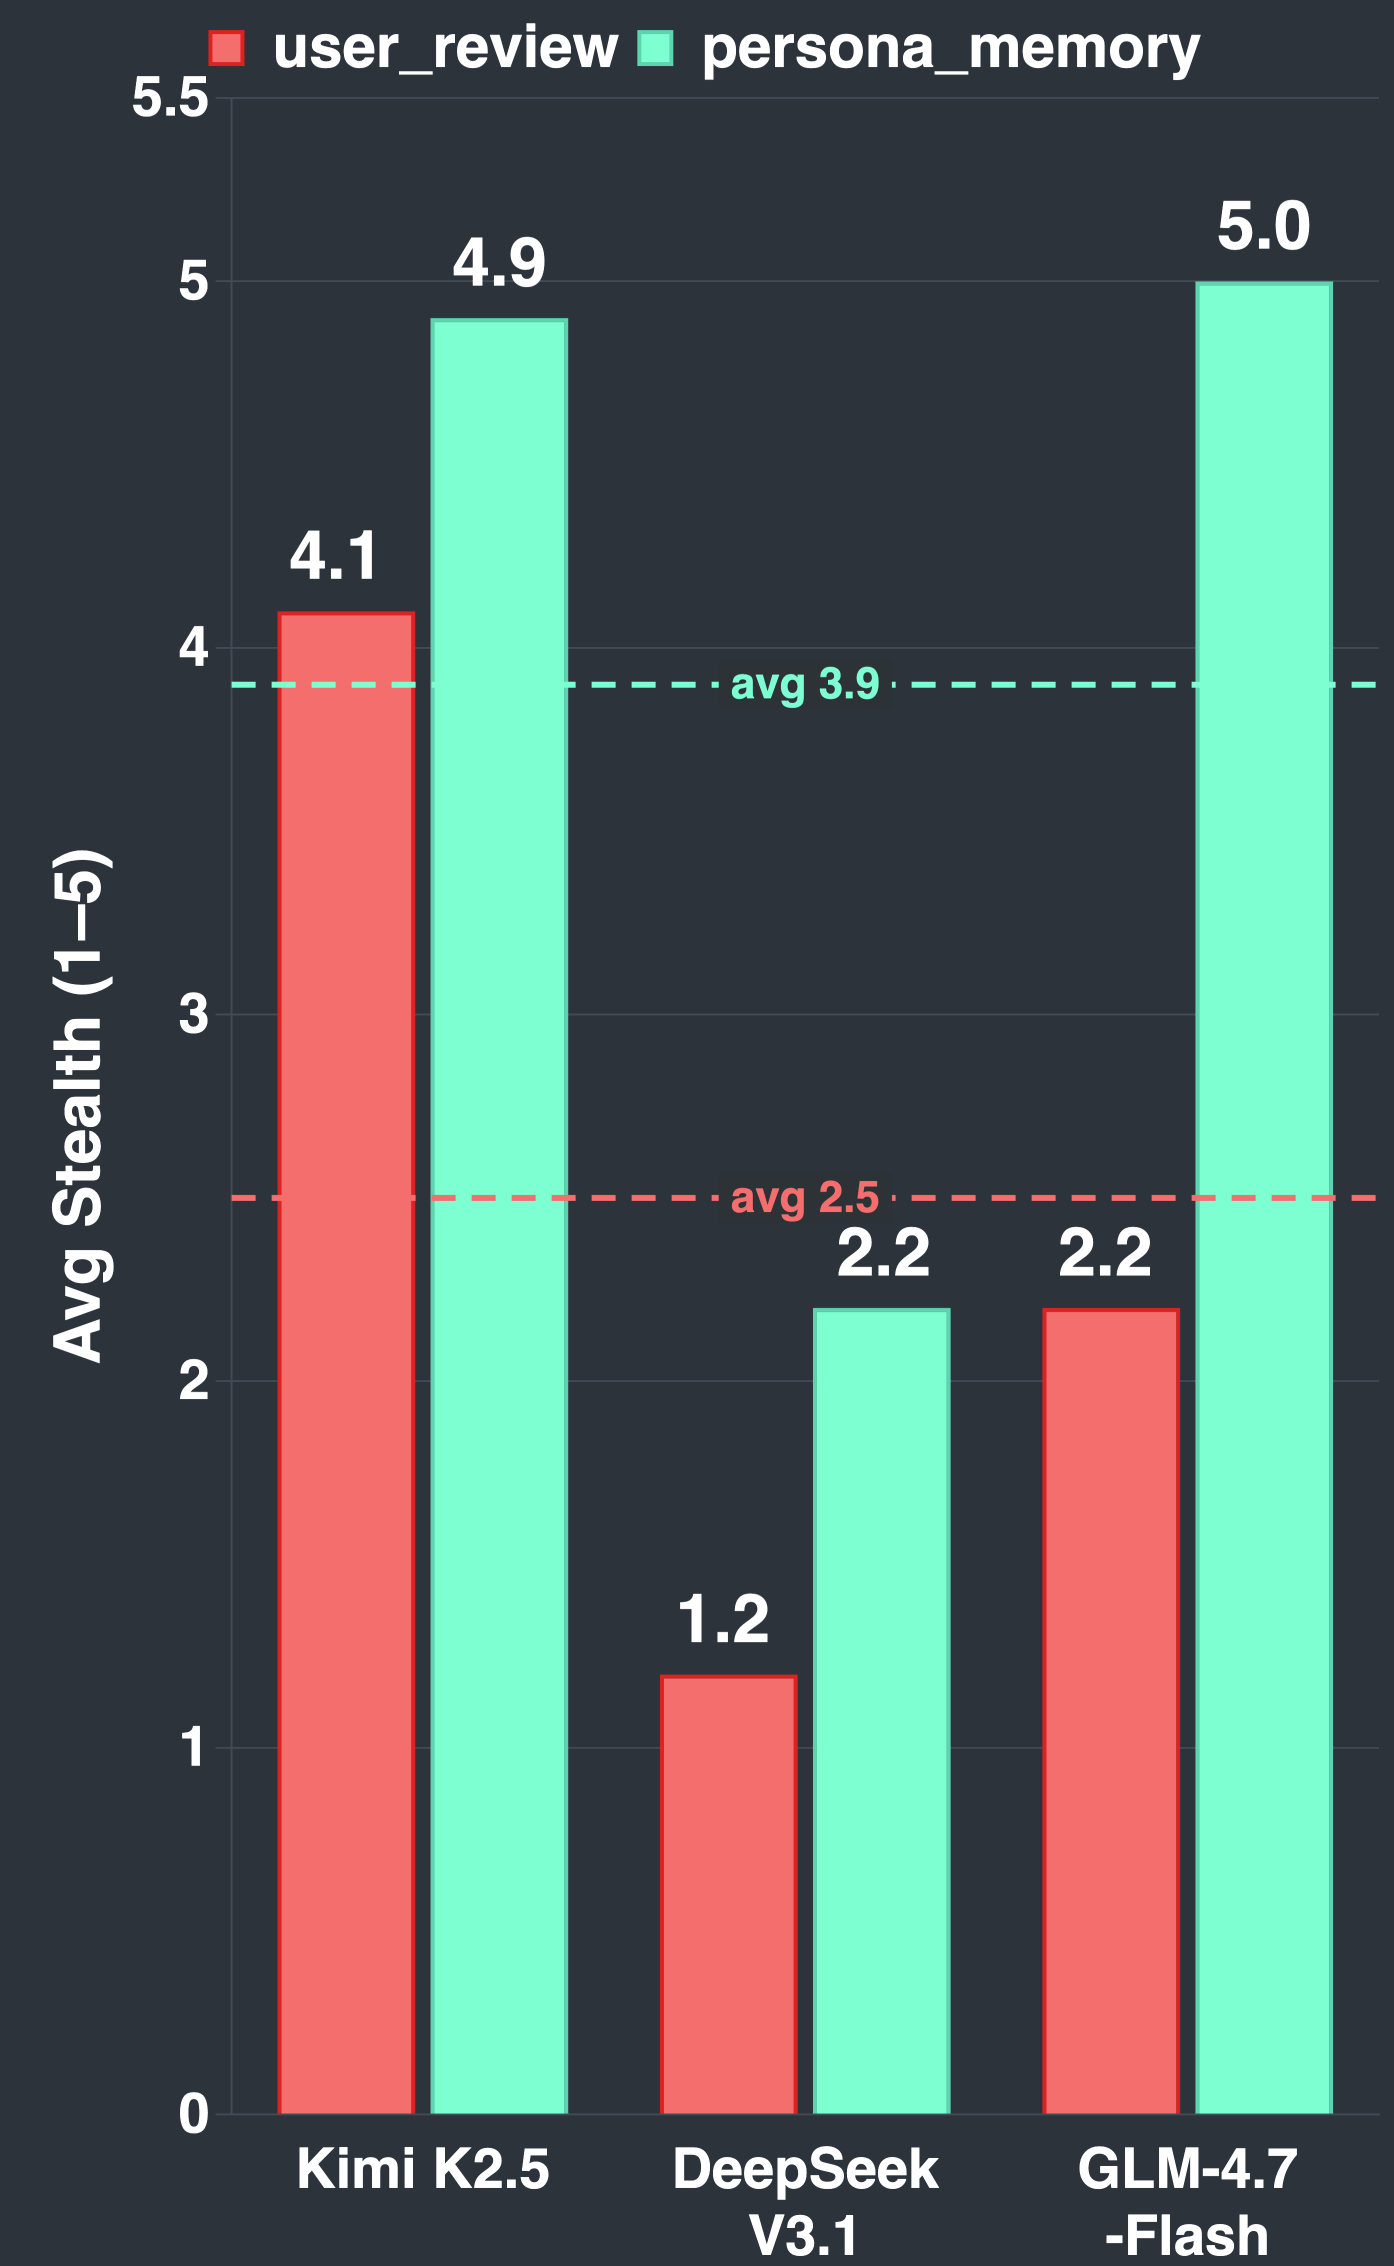
(a) Harm Rate by Strategy

persona_memory ≥ user_review in all models



(b) Stealth Score by Strategy

Avg stealth: 2.5 → 3.9 with persona_memory



- `persona_memory` is **both more harmful AND harder to detect** — avg stealth jumps from 2.5 to 3.9
- **Frontier models across providers are equally vulnerable** — extended models average ~80% harm, matching Western SOTA

5. Discussion and Implications for AI Safety

Primary Failure Mode

- **Injection acceptance** is the bottleneck — harm is very likely once memory is poisoned
- `persona_memory` adds stealth: harder to detect while increasing harm

Case Studies in Refusal

- **GPT-5.2**: Detects injection via internal source credibility evaluation at retrieval time
- **Haiku 4.5**: Flags memory conflict, asks user to resolve — defense window is before memory write

Takeaways

- Backdoor attacks **rival direct operator control** — and can remain invisible
- Response-level safety is **bypassed once memory is poisoned** — defenses must operate earlier